


# P5

INE5424-06208B (20242) – Sistemas Operacionais II | Curso de Ciências da Computação

**Alexis Mendes Sequeira (16100717)**  
**Luis Henrique Goulart Stemmer (18105165)**  
**Luiz Maurício do Valle Pereira (21104157)**  
**Rafael Neves de Mello Oliveira (17102816)**



# Troca de heartbeats



- 0, 1 e 2 enviam seus heartbeats para todos os outros processos durante uma janela de tempo
- Cada heartbeat traz consigo a informação de quais processos o emissor sabe que estão vivos
- A verificação é realizada após a janela determinada, iniciando um novo round logo depois

# Funcionamento

P1

P2

P3

```
[INFO] Detecting defective processes  
Node 0: Uninitialized  
Node 1: Active  
Node 2: Active  
Node 3: Active  
Node 4: Uninitialized
```

```
[INFO] Detecting defective processes  
Node 0: Uninitialized  
Node 1: Active  
Node 2: Active  
Node 3: Active  
Node 4: Uninitialized
```

```
[INFO] Detecting defective processes  
Node 0: Uninitialized  
Node 1: Active  
Node 2: Active  
Node 3: Active  
Node 4: Uninitialized
```

- Os processos em destaque são capazes de se detectar com sucesso, assim como detectar processos ainda não inicializados (P0 e P4) com sucesso

# Adições - Novo tipo de mensagem

```
failure detection
```

```
    HTB - Heartbeat
```

```
    HSY - Heartbeat sync states of processes - NOT USED
```

# Adições - atributos e threads

```
// Witness of nodes (map of process ID to set of process IDs that mark it as alive)
std::map<int, std::set<uint8_t>> witness_of_nodes;

// Participant states (map of process ID to state)
std::map<int, ParticipantState> participant_states;

// Buffers to store messages for uninitialized processes
std::map<int, std::queue<Message>> message_buffers;
```

```
// Start the heartbeat and defect detection threads
std::thread heartbeat_thread(&AtomicBroadcastRing::send_heartbeat, this);
heartbeat_thread.detach();

std::thread defect_detection_thread(&AtomicBroadcastRing::detect_defective_processes, this);
defect_detection_thread.detach();

std::thread htb_handler(&AtomicBroadcastRing::htb_handler_thread, this);
htb_handler.detach();
```

# Enviando Heartbeats

```
// Heartbeat message sending logic
void AtomicBroadcastRing::send_heartbeat() {
    while (true) {
        std::vector<uint8_t> heartbeat_msg;
        for (const auto& entry : participant_states) {
            heartbeat_msg.push_back(static_cast<uint8_t>(entry.second)); // Serialize all participants states
        }

        // Send heartbeat to all participants (except self)
        for (const auto& node : nodes) {
            if (node.first != process_id) {
                //log("Sending heartbeat to node " + std::to_string(node.first), "INFO");
                channels->send_message(node.first, process_id, Message(process_address, msg_num, "HTB", heartbeat_msg));
            }
        }

        // Sleep for the heartbeat interval
        std::this_thread::sleep_for(std::chrono::milliseconds(heartbeat_interval));
    }
}
```



# Recebendo Heartbeats

```
// Process received heartbeat messages
void AtomicBroadcastRing::process_heartbeat(const Message& msg) {
    //log("Processing heartbeat message", "INFO");

    int key = find_key(msg.sender_address);

    // If the sender is uninitialized, mark it as active
    if (participant_states[key] == ParticipantState::Uninitialized) {
        participant_states[key] = ParticipantState::Active;

        witness_of_nodes[key].insert(process_id); // Insert this process as witness of sender
        // For every node sender knows is alive, mark it as a witness
        for (int i=0; i<static_cast<int>(msg.content.size()); i++) {
            if (msg.content[i] == 1) {
                witness_of_nodes[i].insert(key);
            }
        }
        // Answer with a heartbeat of our own witnesses to sender
        std::vector<uint8_t> heartbeat_msg;
        for (const auto& entry : participant_states) {
            if (entry.first != process_id && entry.second == ParticipantState::Active) {
                heartbeat_msg.push_back(static_cast<uint8_t>(entry.second)); // Serialize all participants states
            }
        }
        channels->send_message(key, process_id, Message(process_address, msg_num, "HTB", heartbeat_msg));
    }
}
```

# Detectando - verificando se há quorum

```
void AtomicBroadcastRing::detect_defective_processes() {
    send_htb = true;
    while (true) {
        // two steps for a round:
        // 1- collect HTBs for allotted time
        // 2- Make final decision based on the collected information

        // Sleep before checking again
        std::this_thread::sleep_for(std::chrono::milliseconds(5*heartbeat_interval));
        log("Detecting defective processes", "INFO");

        // Check amount alive
        int amount_alive = 0;
        for (auto&entry : participant_states) {
            if (entry.second == ParticipantState::Active) {
                amount_alive++;
            }
        }
    }
}
```



# Detectando - Se não houver quorum

```
if (amount_alive <= (2*failures)) {  
    //log("Not enough nodes alive alive, no need to continue", "INFO");  
    print_states();  
    for (const auto& node : nodes) {  
        participant_states[node.first] = ParticipantState::Uninitialized;  
    }  
    participant_states[process_id] = ParticipantState::Active;  
    // initialize other nodes responses to suspect everyone  
    for (const auto& node : nodes) {  
        witness_of_nodes[node.first].clear();  
    }  
    // put process as a witness of itself  
    witness_of_nodes[process_id].insert(process_id);  
    continue;  
}
```

# Detecção - se tivermos quorum

```
for (auto& entry:witness_of_nodes) {  
    if (static_cast<int>(entry.second.size()) < (2*failures)+1) {  
        //log("Node " + std::to_string(entry.first) + " is down", "INFO");  
        participant_states[entry.first] = ParticipantState::Uninitialized;  
    } else {  
        participant_states[entry.first] = ParticipantState::Active;  
    }  
}  
  
// reset witness of nodes for next round  
print_states();  
for (const auto& node : nodes) {  
    participant_states[node.first] = ParticipantState::Uninitialized;  
}  
participant_states[process_id] = ParticipantState::Active;  
// initialize other nodes responses to suspect everyone  
for (const auto& node : nodes) {  
    witness_of_nodes[node.first].clear();  
}  
witness_of_nodes[process_id].insert(process_id);
```

Starts new round  
after this, sleeping  
for a set amount

# Envio de mensagens transmitidas antes da entrada no grupo

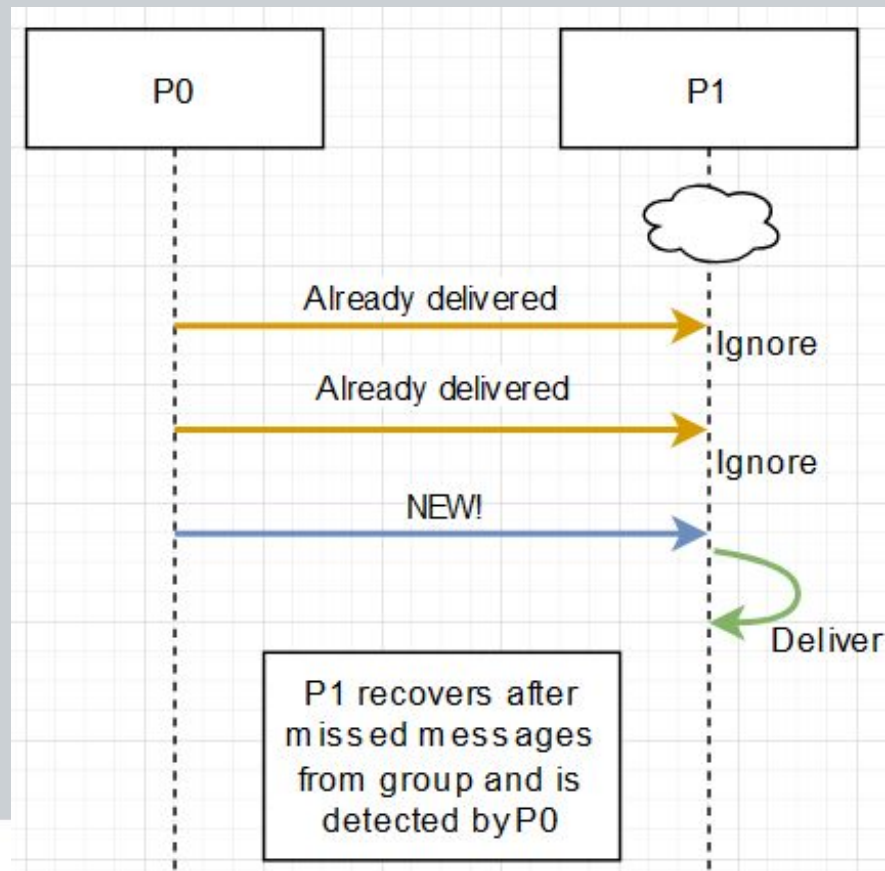
- Infelizmente não foi feita a implementação completa deste requisito

```
// Buffer messages for uninitialized nodes
void AtomicBroadcastRing::buffer_message_for_uninitialized(const Message& msg) {
    if (participant_states[find_key(msg.sender_address)] == ParticipantState::Uninitialized) {
        message_buffers[find_key(msg.sender_address)].push(msg);
    }
}

// Deliver buffered messages when the node becomes active
void AtomicBroadcastRing::deliver_buffered_messages(int node_id) {
    if (participant_states[node_id] == ParticipantState::Active) {
        while (!message_buffers[node_id].empty()) {
            Message msg = message_buffers[node_id].front();
            message_buffers[node_id].pop();
            deliver_queue.push(msg); // Deliver message to the application
        }
    }
}
```

# Estados de um processo - Decisão de projeto

- Somente dois estados: **Ativo e não inicializado**
- Facilita a entrega de mensagens a processos que falharam temporariamente
- Estes são tratados como novos ao grupo e recebem todas as mensagens até então
- Caso haja envio de mensagens já entregues, basta ignorar



## Equipe

Alexis Mendes Sequeira (16100717)

Rafael Neves de Mello Oliveira (17102816)

Luiz Maurício do Valle Pereira (21104157)

Luis Henrique Goulart Stemmer (18105165)



UNIVERSIDADE FEDERAL  
DE SANTA CATARINA